

The Anthropic Incident Reveals What AI Governance Is Missing

A Civilian Certification Body for AI: By the People, For the People

Mike Rizek

Author and framework developer — constitutional theory, geopolitical security, AI governance

June 15, 2026

On Friday, June 12, 2026, the US government issued an export control directive requiring Anthropic to suspend all access to its two most capable AI models — Fable 5 and Mythos 5 — by any foreign national, whether inside or outside the United States, including the company's own foreign national employees. Unable to verify nationality in real time, Anthropic disabled both models for every customer worldwide.

The directive arrived at 5:21pm Eastern. By the company's own account, the letter did not provide specific details of its national security concern. The government's understanding, as Anthropic described it, was that a method had been found of bypassing safeguards intended to limit the model's use in identifying software vulnerabilities. Anthropic called the issue narrow and said it had not received detailed written support for the government's position.

That is the state of AI governance today.

What is documented is enough to make the point. A claim about a vulnerability reached the government through channels that were never made public. A directive followed within an unknown interval, citing national security authorities and providing no specific detail. Access was disabled globally within hours. Between the claim and the action, there is no visible step. No independent technical review. No proportionality assessment. No institutional mediation between a contested claim and a consequential order.

This is not a failure of individuals. It is a failure of architecture.

The Problem Is Structural

Governments are not equipped to lead AI governance. Political cycles reset every two to four years. Institutional memory evaporates with each administration. The technical education that fills that vacuum comes largely from lobbyists and commercially interested parties — companies framing their competitors as dangers, executives seeking headlines, actors signaling loyalty to administrations for access and advantage.

As AI becomes more deeply integrated into healthcare, financial systems, legal processes, and public infrastructure, the cost of this structural failure rises faster than the technology itself. A rapid government response to a contested vulnerability claim is no longer merely irresponsible. In a sufficiently integrated AI environment, it becomes a threat vector in its own right. The ability to trigger a government overreaction is the attack surface.

The current system has no filter for that. Everything arrives at government with equal urgency, regardless of merit, regardless of motive.

What Is Actually Needed

AI public infrastructure requires a civilian certification body — independent, constitutionally grounded, and accountable to the people rather than to government or industry.

Not a regulatory agency. Not an industry consortium. Not an advisory committee. A civil authority whose mandate derives from public interest, whose findings government receives rather than directs, and whose certification is required before any AI system — from any developer — reaches the public.

The funding model. All AI companies collectively finance the body through obligatory contribution. No single company's contribution is large enough to purchase influence. Withdrawal is not an option. This is the FINRA model applied to AI: industry bears the cost of oversight without controlling its direction.

The membership. Drawn from existing bodies with demonstrated technical credibility and institutional independence — FINRA, IAEA, ICAO, IEEE — supplemented by technical government agency representatives contributing expertise without political direction, and an independent judicial council, structurally separate from the Supreme Court and its appointment politics but carrying recognized standing respected by courts, ensuring proportionality, due process, and constitutional reasoning are embedded in every evaluation.

The government role. Audit and receipt. Government verifies the body's process and financial integrity. It acts on certified findings. It does not appoint, direct, or control. The direction of accountability runs from civil society upward to government, not downward from government to the committee.

The single channel principle. All substantive AI communications to government pass through the committee. No bilateral lobbying. No vulnerability claims entering political channels without independent technical evaluation. One trusted channel, trusted because it is the only channel.

The Deployment Threshold

Internal research remains free. Companies develop, experiment, and compete without restriction at the research level. But the moment a system crosses from internal development to public deployment, it enters the committee's jurisdiction.

That threshold — research to public implementation — is where certification activates. It is where capability meets consequence. It is the appropriate intervention point.

This applies without exception. A government-developed AI system intended for public use goes through the same certification process as a commercial model. Government cannot be developer, deployer, and certifier simultaneously. That concentration of roles is precisely the condition that produces unchecked power.

Sovereign AI development for strategic and military use operates in a separate domain, as it always has. Governments will compete in AI as they compete in every strategic technology. That competition belongs to governments, not to the public. What the public needs is the opposite of what arms races produce: consistency, standards, reliability, and trust across AI infrastructure that affects their lives daily.

Governments compete. The public deserves consistency. The committee is the institutional boundary between those two domains.

From Fragmentation to Coherence

The current AI governance landscape is fragmented by design and by accident. Companies developing independently. Governments responding reactively. International bodies without enforcement authority. Standards organizations without deployment power. Professional communities talking past political ones.

Every major AI incident is partly a fragmentation failure.

What a civilian certification body creates is not centralized control. It creates coherence: a single institutional center where fragmented expertise, civil concern, government interest, and industry knowledge converge into a unified evaluative process with legitimate authority.

This is the next step in building AI trust infrastructure. Not conditions alone. Not framework alone. The institutional form that makes constitutional principles operational at the scale the technology now demands.

By the people. For the people. Without exception.

Postscript, September 2026

On June 30, 2026, the export controls on Fable 5 and Mythos 5 were lifted, and access was restored on July 1. The suspension lasted eighteen days.

The resolution does not change the argument. It confirms it. What ended the suspension was negotiation between a government and a company, conducted out of public view, on a timeline neither the affected public nor the affected industry could observe or influence. The same absence that produced the order produced its reversal. Eighteen days of global disruption passed through no institution built to evaluate the claim, weigh the proportionality of the response, or represent the hundreds of millions of people the decision reached.

An architecture that can be switched off by a letter and switched back on by a phone call is not governed. It is administered.

This proposal extends the constitutional framework developed in Trust Architecture: A Constitutional Framework for AI Infrastructure. The author welcomes engagement from governance researchers, technologists, legal scholars, and institutional designers working on these questions.